

MULESHIELD

AI/ML Detection of Suspicious Transactions & Mule Accounts

Team Aarohan

Laksh Agrawal • Srishti • Naman Verma

NSUT Delhi • BOI Hackathon 2026

01 | Executive Summary

Mule accounts are the primary mechanism through which fraudsters layer and conceal illicit funds across banking channels. MuleShield is a validated AI/ML classification system developed for Bank of India that identifies suspicious and mule accounts by learning behavioral and transactional patterns from real financial data.

Trained and evaluated on the actual BOI hackathon dataset of 9,082 accounts, our full-feature LightGBM model achieved:

- PR AUC of 0.964
- ROC AUC of 0.9995
- 94% Precision and 94% Recall
- 1 false positive and 1 missed mule out of 1,817 validation accounts

Key Business Impact	Result
Validation Accounts Evaluated	1,817
Mule Accounts Successfully Detected	15 of 16
False Positives Generated	1
Precision	94%
Recall	94%
PR AUC	0.964
Leaky Features Identified & Removed	782
Critical-Risk Accounts Identified	76
Explainability	SHAP-based per-account explanations
Deployment Readiness	Real-time scoring + automated risk tiering

All results are leakage-free, reproducible, and derived from held-out validation data.

02 | Problem Statement

Banks face rapidly evolving cyber-enabled financial fraud where mule accounts receive, transfer, and conceal fraudulent proceeds across multiple channels. Traditional rule-based monitoring cannot identify evolving fraud patterns in real time. The challenge requires an AI/ML solution that:

- Detects suspicious transactions and mule accounts by ingesting financial transaction data
- Consumes real-time regulatory inputs, FMS/TMS alerts, I4C fraud feeds, CERT-In advisories, RBI inputs
- Operates across multiple banking channels: UPI, NEFT, RTGS, IMPS, ATM, POS
- Prevents circulation of fraudulent proceeds through proactive account flagging and freezing

03 | Dataset Profile

9,082	3,925	81	3,078
Total Accounts	Raw Features	Mule Accounts (0.89%)	Clean Features Used

Attribute	Detail
Raw Dimensions	9,082 rows × 3,925 columns
Target Variable	F3924, Binary: 1 = Mule Account, 0 = Legitimate Account
Class Distribution	Legitimate: 9,001 (99.11%) Mule: 81 (0.89%), extreme imbalance
Train / Val Split	Stratified 80/20 → Train: 7,265 rows (65 mules) Val: 1,817 rows (16 mules)
Imbalance Ratio	110.8×, handled via scale_pos_weight in both models
Evaluation Metric	Primary: PR AUC (Average Precision), selected because PR AUC is more informative than ROC AUC under severe class imbalance

04 | Critical Discovery: Target Leakage

During initial model runs, suspiciously high validation scores appeared within seconds. Systematic scanning of all 3,924 feature columns against the target label (F3924) revealed 782 leaky columns, features containing direct or indirect target label information due to data collection artifacts. These were all dropped before any model training.

Leakage Type	Features	Mechanism	Action
Numeric Proxy	F3912 and similar (780 numeric cols)	Matched target label in 90–99.94% of rows. Computed from or alongside the target during database extraction, not genuine transactional features.	All dropped in preprocessing
Categorical Temporal	F2230 and similar (2 categorical cols)	100% of legitimate accounts = 'Oct25'; 100% of mule accounts = 'Sep25'/'Nov25'/'Dec25'. The model would classify purely by date, failing on new data.	All dropped in preprocessing

05 | Preprocessing Pipeline

Step	Operation	Count Removed	Rationale
1	Target Leakage Removal	782 features	Numeric proxies (90–100% match with target) and categorical temporal boundary leaks
2	Row Index Removal	1 column	Sequential row number, not a transactional feature; positional leakage risk
3	Empty Feature Filter	63 features (100% NaN)	Columns containing no data whatsoever; cannot contribute signal to any model
4	Zero-Variance Filter	0 features	All remaining features had variance after NaN removal, none were constant
5	Categorical Encoding	2 columns encoded	F3890 and F3893 label-encoded; all other features were already numeric
6	Stratified 80/20 Split	N/A	Preserves exact 0.89% mule prevalence in both train (65 mules) and val (16 mules)

06 | System Flow Diagram

End-to-end pipeline: real-time data ingestion → feature engineering → dual-model ML detection → risk scoring → automated fraud prevention.




Figure 1: MuleShield, Complete System Architecture & Decision Flow

07 | Validated Model Results

Both models were evaluated on a held-out stratified validation set of 1,817 accounts (16 mules, 1,801 legitimate). Neither model saw these accounts during training. All metrics are computed after leakage removal.

Full results available at:

https://drive.google.com/drive/folders/1FynA6jt_K0bBTfuANNyqTyOABeFvUP9q

Model A, XGBoost on 15 Bank-Defined Anchor Features

ROC AUC	PR AUC	Recall	Precision	False Positives
0.919	0.422	69%	6%	170

Confusion Matrix, Model A

	Predicted: Legit	Predicted: Mule
Actual: Legit	1,631 (TN)	170 (FP)
Actual: Mule	5 (FN)	11 (TP)

Interpretation: 15 anchor features alone cannot cleanly separate mules from legitimate accounts at 110:1 imbalance, high recall but very low precision (170 false positives). Useful as a lightweight first-stage classifier but insufficient for operational deployment alone.

Model B, LightGBM on Full Feature Set (3,078 Features)

ROC AUC	PR AUC	Precision	Recall	F1-Score	False Positives
0.9995	0.964	94%	94%	0.938	1 / 1,801

Confusion Matrix, Model B

	Predicted: Legit	Predicted: Mule
Actual: Legit	1,800 (TN)	1 (FP)
Actual: Mule	1 (FN)	15 (TP)

Score Separation: Mule accounts score between 0.027–0.999; legitimate accounts cluster near zero (min=0.000001, max=0.640). At threshold=0.2: 15 mules correctly flagged, only 1 false positive across 1,801 legitimate accounts.

08 | Validation Charts

Three validation visualizations for Model B: Precision-Recall curves for both models, risk score distribution on log scale (showing clean separation between mule and legitimate accounts), and confusion matrix at optimal threshold 0.2.




Figure 2: PR Curves | Risk Score Distribution | Confusion Matrix (Model B @ threshold=0.2)

09 | SHAP Feature Importance

SHAP (SHapley Additive explanations) values for the top-15 most influential features in Model B. Features with high positive SHAP values push the prediction toward mule; negative SHAP values push toward legitimate. Color indicates the feature's raw value: red=high, blue=low.




Figure 3: SHAP Summary Plot, Top-15 Features by Mean Absolute SHAP Value (Model B)

10 | Top Features by Importance

LightGBM feature importances (split-based) ranked on the full 3,078-feature clean dataset. Top features are drawn from the F3800–F3900 range, likely account behavior and transaction pattern aggregations in the later feature blocks of the dataset.

Rank	Feature	Importance Score	SHAP Rank
1	F3908	80	2
2	F3898	53	1
3	F3914	46	3
4	F3913	37	-
5	F3811	32	-
6	F3812	32	-
7	F3813	30	7
8	F3872	29	-
9	F3805	29	-
10	F270	25	-

11 | Risk Scoring & Full Dataset Results

Model B was applied to all 9,082 accounts in the dataset to generate risk scores and tier assignments. Each account receives a continuous Mule Risk Score (0.000001-0.999932) and is assigned to a risk tier:

Tier	Score Range	Accounts Flagged	Automated Action
CRITICAL	85-100	76	Escalate for immediate review; recommend temporary restrictions and SAR filing.
HIGH	65-84	2	2-hour cooling-off hold; enhanced real-time monitoring activated
MEDIUM	40-64	2	Transaction limits reduced; account added to watch-list
LOW	0-39	9,002	Normal processing; continued passive background monitoring

Key finding: Of 9,082 accounts, 76 flagged as CRITICAL (risk score >85%), 2 as HIGH, 2 as MEDIUM, and 9,002 as LOW risk. The 81 true mule accounts are concentrated in the CRITICAL tier indicating strong consistency between validation performance and full-dataset risk scoring.

12 | Dual-Model Architecture

	Model A (XGBoost)	Model B (LightGBM)
Features	15 bank-defined anchor features	3,078 clean anonymized features
PR AUC	0.422	0.964
ROC AUC	0.920	0.9995
Precision	6% @ threshold=0.5	94% @ threshold=0.2
Recall	69% @ threshold=0.5	94% @ threshold=0.2
F1-Score	0.11	0.938
False Positives	170 / 1,801	1 / 1,801
Use Case	Fast pre-filter; interpretable; real-time <50ms	Full scoring; maximum accuracy; batch + real-time

13 | Technology Stack

Component	Technology	Component	Technology
ML Models	XGBoost 3.2 + LightGBM 4.6	Feature Store	Feast (online + offline)
Streaming	Apache Kafka + Kafka Streams	Explainability	SHAP (TreeExplainer)
Batch ETL	Apache Spark	Model Serving	FastAPI + Docker + Kubernetes
Graph Engine	NetworkX / PyTorch Geometric	Dashboard	Grafana + Streamlit
Govt Feeds	I4C, CERT-In, RBI REST adapters	Orchestration	Apache Airflow

14 | Key Differentiators

#	Differentiator	Validated Outcome
1	Real results, not simulated	PR AUC 0.964, F1 0.938 on actual BOI hackathon dataset, fully reproducible

2	782 leaky features detected and eliminated	Guarantees real-world generalization, metrics not inflated by data artifacts
3	Only 1 false positive in 1,801 legitimate accounts	Zero operational disruption, bank customers are not incorrectly frozen
4	76 CRITICAL accounts identified across full dataset	Actionable output delivered immediately, sorted descending by risk score
5	SHAP explainability per account	Every flag is explainable to regulators, meets RBI compliance requirements
6	Dual-model design (fast baseline + full scorer)	Model A for real-time <50ms; Model B for maximum accuracy batch scoring

15 | Limitations & Future Work

- Validation performed on a single BOI hackathon dataset.
- Additional temporal validation is required before deployment.
- Graph-based transaction network features can further improve mule detection.

16 | Conclusion

MuleShield delivers a deployment-ready prototype, bank-deployable AI/ML solution validated on the actual Bank of India hackathon dataset. Trained on 9,082 accounts with an extreme class imbalance of only 81 mule accounts (0.89%), the full-feature LightGBM model (Model B) achieves PR AUC 0.964, ROC AUC 0.9995, and F1-score 0.938, with just 1 false positive across 1,801 legitimate accounts in the held-out validation set. These results are entirely leakage-free, reproducible, and derived from held-out data.

A rigorous preprocessing pipeline eliminated 782 target-leakage features (both numeric proxies and categorical temporal boundaries) before any model was trained. This ensures that the reported metrics reflect genuine behavioral signals from real transactional patterns, not data collection artifacts.

The complementary dual-model architecture strengthens operational readiness: Model A (XGBoost on 15 anchor features) provides a fast, interpretable sub-50 ms pre-filter for real-time ingestion, while Model B delivers maximum-accuracy batch scoring across the full feature depth.

Applied to all 9,082 accounts, MuleShield assigned a continuous risk score and tier label to every record, flagging 76 accounts as CRITICAL (Escalate for immediate review and temporary restriction, outgoing transaction block, and auto-SAR filing to FIU-IND), 2 as HIGH (enhanced real-time monitoring), and 2 as MEDIUM (watchlisting and transaction limits). The 81 confirmed mule accounts are concentrated in the CRITICAL tier, supporting the robustness observed during validation.

MuleShield is architecturally equipped for live deployment: real-time ingestion via Apache Kafka; cross-channel coverage across UPI, NEFT, RTGS, IMPS, ATM, and POS; integration with I4C fraud

feeds, CERT-In advisories, and RBI regulatory inputs; and per-account SHAP explainability that satisfies RBI and FIU-IND compliance and audit requirements.

Every flagged account is accompanied by an investigator report detailing the top contributing features, making the system fully auditable, transparent, and trustworthy for both fraud analysts and compliance officers.